

Linking urban social innovation to political outcomes: a research design memo

The dataset this project needs does not yet exist—but the building blocks are closer to assembly than ever. No cross-national, city-level database currently integrates social innovation indicators with electoral outcome data. Yet a convergence of new resources—the American Local Government Elections Database (2023), expanding CDP climate disclosures (~1,000 cities globally), the LocalView corpus of 2,861 U.S. local government transcripts, ([Nature](#)) and forthcoming European municipal governance data (Lo-QoG)—makes a staged construction strategy viable for the first time. This memo maps the data landscape honestly, identifies where global comparison is feasible and where within-country analysis is the realistic fallback, and proposes a modular research design. The core intellectual contribution this project can make is constructing the first **city-level analogue of the Walker/Boehmke & Skinner state policy innovativeness score** and linking it to electoral consequences—a question that, remarkably, the literature has not yet answered directly.

The conceptual gap this project fills

The team's framework draws on two mature but disconnected literatures. Urban sociology, through Brandtner (2022 AJS; 2024 Nature Cities), has established that social innovation adoption in cities is *distributed* across public and private organizations and driven by civic capacity—the density of values-oriented nonprofits that recognize and act on social problems. ([DOI Resolver](#)) ([University of Chicago Press](#)) Political science, through Boehmke & Skinner (2012), Caughey & Warshaw (2022), and Shipan & Volden (2008, 2012), has built sophisticated infrastructure for measuring policy innovativeness and diffusion at the *state* level. **The missing link is applying innovativeness measurement at the city level and connecting it to electoral outcomes.**

The SPID database (Boehmke et al., 2020) tracks adoption timing for ~730 policies across U.S. states and computes dynamic innovativeness scores using event-history analysis. ([SSRN +2](#)) It is the gold standard. No city-level equivalent exists at comparable scale. The URBINSOC framework's ambition—operationalizing "urban social innovativeness" as a measurable city characteristic following Walker (1969)—would require building a database of **50–100 trackable policy innovations** across hundreds of cities, recording adoption dates, and computing analogous scores. This is ambitious but achievable with current data infrastructure.

Brandtner's recent trajectory confirms the feasibility. His 2024 Nature Cities paper surveyed 863 civil society organizations across 536 neighborhoods in five global cities (San Francisco, Seattle, Shenzhen, Sydney, Vienna), demonstrating cross-national organizational measurement at the neighborhood level. ([DOI Resolver +2](#)) His 2025 ASR paper with Sharkey, Young, and Bergemann traced the diffusion of opioid litigation across U.S. local governments using event-history analysis— ([Sage Journals](#)) precisely the method needed for city-level innovativeness scoring. The analytical toolkit exists; the bottleneck is the underlying policy adoption data.

What the current database landscape actually offers

The data environment is **asymmetric in three dimensions**: climate policy data vastly outstrips social innovation data; U.S. coverage dwarfs non-U.S. coverage; and election data is more available than policy

adoption data at the city level.

For climate policy, CDP-ICLEI Track provides the most comprehensive global resource: roughly 1,000 cities annually reporting CDP greenhouse gas emissions, adaptation actions, mitigation targets, and climate finance pipelines CDP through an open data portal with structured, downloadable datasets. CDP CDP The Global Covenant of Mayors adds breadth (13,500+ signatories) but with variable reporting quality.

Globalcovenantofmayors Globalcovenantofmayors These databases are mature, standardized, and API-accessible—the strongest foundation for a global comparative component of the project, albeit limited to environmental policy.

For broader U.S. policy adoption, several resources can be assembled. CityHealth, which the team may not know about, tracks adoption of 12 evidence-based health policies across the 75 largest U.S. cities City Health annually since 2017, rated on a gold/silver/bronze scale—an excellent structured panel dataset for statistical analysis. ICMA surveys cover ~11,000 municipalities on management practices, service delivery, and institutional structure, ICMA with waves going back to 1982. ICMA ICMA The NLIHC Tenant Protections Database (launched 2021) systematically tracks local housing policy innovations including eviction protections, right-to-counsel laws, and rent stabilization. National Low Income Housing ... Bloomberg's What Works Cities Certification now covers 104+ cities with 43 criteria measuring data-driven governance capacity, City of Philadelphia and critically expanded to Canada and Latin America in 2024–2025. What Works Cities +2

For text-as-data approaches, LocalView provides 139,616 transcribed local government meeting videos from 2,861 governments (2006–2022), freely available and linkable via FIPS codes to Census and election data. Nature ResearchGate This enables NLP-based measurement of policy attention, agenda-setting speed, and innovation diffusion from textual evidence—but requires substantial researcher effort to extract structured innovation measures from raw transcripts. The text-reuse methodology that Jansa, Hansen, and Gray (2019) applied to state legislatures has not yet been adapted to municipal ordinances at scale, though the approach is directly transferable if systematic ordinance corpora can be assembled from sources like Municode or American Legal Publishing.

For organizational infrastructure underlying innovation, the IRS/NCCS nonprofit data provides comprehensive annual coverage of U.S. nonprofits by geography and type, enabling the civic capacity measures Brandtner validated. The Democracy at Work Institute tracks worker cooperatives; the Participatory Budgeting Project tracks PB adoption across U.S. and Canadian cities; Participatory Budgeting Project and the PB Atlas (pbatlas.net) documents over 11,500 participatory budgeting processes globally as of 2024. Wikipedia

For European cities, the landscape is improving but fragmented. The European Social Innovation Database (ESID) contains 11,468 projects from 159 countries, classified by machine learning with F1 scores up to 0.90. PubMed Nature However, these are project-level records requiring geocoding and aggregation to the city level. The European Social Enterprise Monitor (ESEM), now in its third edition, covers 30 countries and 1,807 social enterprises—Euclidnetwork growing but still enterprise-level, not city-level. The Decidim participatory governance platform, European Institute of Public Ad... used by 300+ instances across 20+ countries, generates structured participation data but without centralized aggregation. Wikipedia The European Urban Initiative (successor to Urban Innovative Actions) funds innovative urban projects with €450 million in ERDF resources for 2021–2027, Urban-initiative Urbanagendaplatform providing a database of funded innovation projects searchable by city. Europa Europa

For Asian and Latin American cities, the picture is stark. No structured, quantitative, city-level policy adoption database exists for Asian cities beyond CDP's climate reporting. Latin America has stronger foundations: Brazil's TSE provides comprehensive municipal election data widely used in econometric research, and Bloomberg's What Works Cities expansion now includes cities like Bogotá, Medellín, Porto Alegre, and Recife. (What Works Cities) But no LAC-wide equivalent of CityHealth or ICMA exists. ECLAC's social innovation database documented 4,800 experiences across competition cycles but is not an ongoing tracking system. (CEPAL)

What recent political science reveals about the core question

The project's central hypothesis—that urban social innovation affects electoral outcomes, potentially including backlash against incumbents—sits in a **well-theorized but empirically underpopulated space**. A systematic review of 2020–2025 literature reveals that the direct link from municipal policy adoption to incumbent electoral punishment has not been rigorously tested at scale.

The closest work comes from several directions. De Benedictis-Kessner and Warshaw (2020, APSR) demonstrated that voters hold local officials accountable for economic conditions but suffer from "attribution confusion," simultaneously rewarding or punishing local officials for national and state economic trends. This finding is crucial: if voters struggle to attribute economic outcomes correctly, they may also struggle to attribute policy innovations to the responsible officials, weakening both reward and punishment mechanisms. Stokes's backlash framework, extended in her 2023 PNAS work on wind energy opposition, establishes that spatially concentrated policy costs (like wind turbines within 3 km) can reduce incumbent vote shares by 4–10%—(IDEAS/RePEc) but this operates at the project-siting level, not the policy-adoption level. (Utxas)

On policy feedback, Béland, Campbell, and Weaver's 2022 Cambridge Elements volume explicitly notes that **most policy feedback research focuses on national programs**; local-level feedback remains theoretically underdeveloped. (The Roosevelt Institute) Michener's (2023) analysis of pandemic-era policies found "fraught political paths and apparent absence of short-term positive feedback"—(The Roosevelt Institute) a cautionary finding for expecting clean feedback loops from municipal innovation. An emerging framework from Fullerton and colleagues treats **state preemption as a form of negative policy feedback**, arguing that when states preempt local innovations (in labor, housing, civil rights), this produces both resource effects (reducing local capacity) and interpretive effects (signaling whose rights matter). (National League of Cities +2)

The BLM/police reform episode provides the most vivid recent case study. A 2024 Social Problems paper analyzing hand-compiled budget data from 264 major U.S. cities found no evidence that protests led to police defunding; in cities with large Republican vote shares, protests were associated with *increased* police budgets. (ResearchGate) The Minneapolis Question 2 vote in 2021, which proposed replacing the police department, was soundly defeated even in a strongly liberal city—(Taylor & Francis Online) demonstrating that progressive policy preferences are more nuanced than partisan affiliation suggests. (Taylor & Francis Online) Bolet, Green, and González-Eguino (2023, APSR) found that Spain's "just transition" agreements for coal communities could mitigate electoral backlash against climate policy, suggesting that policy design matters enormously for political outcomes. (University of Chicago Press)

Three structural features of local politics complicate the innovation-backlash link. First, Anzia's (2022) *Local Interests* documents how organized interest groups—public employee unions, developers, business associations

—disproportionately shape local policy, (Goldman School of Public Policy) suggesting that innovation may be more responsive to organized interests than to voters. Second, Einstein et al. (2025) show that **age and homeownership drive enormous local turnout gaps**, with older homeowners exerting disproportionate influence— (Sage Journals) precisely the demographic most likely to resist disruptive innovations.

(National Civic League) Third, Schaffner, Rhodes, and La Raja (2024, PSRM) find that local policy attitudes are structured along the same ideological dimension as national attitudes, (ResearchGate) (University of Chicago Press) challenging the older view that local politics operates on fundamentally different dimensions and suggesting that national partisan dynamics may swamp local innovation effects. (Cambridge Core)

Datasets that enable linking innovation to elections

The critical infrastructure question is whether social innovation data and election data can be joined at the city level. The answer varies dramatically by country.

In the United States, linkage is now feasible at meaningful scale. The American Local Government Elections Database (de Benedictis-Kessner et al., 2023, Scientific Data) provides **~78,000 candidates across ~57,000 electoral contests** in most medium and large cities, with vote shares, candidate partisanship, and demographics. (DNTB +3) This can be merged with ICMA surveys, CityHealth policy data, Census of Governments fiscal data, and NCCS nonprofit data via FIPS codes. (ICMA) The Ferreira and Gyourko mayoral elections dataset (extended through 2017) is specifically designed for regression discontinuity analysis around close elections and has been successfully linked to fiscal, crime, and housing permit data. VEST precinct-level data with GIS boundaries enables spatial linkage of statewide election results to city boundaries. **The U.S. should be the primary empirical setting.**

In Europe, EU-NED (Schraff, Vergioglou & Demirci, 2023) provides harmonized election data for 30 countries at the NUTS 2/3 level from 1990–2020, integrated with PartyFacts party codes and Eurostat regional statistics. NUTS 3 regions approximate municipalities in some countries (especially Scandinavia). Germany offers the strongest municipal-level data: a comprehensive database of municipal, state, and federal election results harmonized to 2021 boundaries. (Princeton University) An 18-country European Municipal Registration and Turnout Database covers local elections from 2013–2020 at the municipality level. (Princeton University) The forthcoming **Lo-QoG** (Local Quality of Government) dataset from the University of Gothenburg, which will use EU Local Administrative Units corresponding to municipalities, could be transformative for European analysis. (Rj) A realistic European strategy would begin with Germany, Scandinavia, and Spain—countries with both strong municipal election data and meaningful local policy variation.

In Latin America, Brazil is the clear first choice. Its TSE provides comprehensive municipal election data that has supported a robust literature including regression discontinuity designs (De Magalhães 2012) and the landmark Ferraz and Finan (2008, 2011 AER) study of audit-driven electoral accountability. Brazil also has participatory budgeting data going back to Porto Alegre in 1989 (Wikipedia) and a growing What Works Cities presence.

Globally, no integrated dataset exists. The CLEA archive covers 170+ countries for national legislative elections at the constituency level (NYU Guides) but not the municipal level. V-Dem provides country-level assessments of local democratic quality, useful for cross-national controls but not within-country city variation.

True global comparison is feasible only for climate policy using CDP data linked to national legislative election results aggregated to urban constituencies—a meaningful but narrow slice of social innovation.

Research designs that have worked and what they teach us

The most successful designs linking municipal policy to political outcomes share a common structure: they exploit **quasi-random variation in who governs** (via close-election RD) or **quasi-random variation in policy exposure** (via spatial or temporal discontinuities), then trace consequences to measurable outcomes.

- **Close-election RD:** De Benedictis-Kessner and Warshaw (2016, 2020) use narrow mayoral election margins to estimate the causal effect of Democratic vs. Republican mayors on fiscal policy, finding real partisan effects on spending and debt. (University of Chicago Press) This design requires the ALGED or Ferreira-Gyourko data and works best with hundreds of elections.
- **Spatial exposure designs:** Stokes (2016) matched geocoded wind turbine locations to polling-station election results in Ontario, exploiting quasi-random variation in turbine placement to estimate incumbent vote losses. (Utxas) (IDEAS/RePEC) This template is directly applicable to spatially concentrated social innovations (e.g., community land trusts, participatory budgeting districts).
- **Policy audit natural experiments:** Ferraz and Finan (2008, 2011) exploited Brazil's random municipal audit program to show that corruption exposure reduces incumbent re-election by substantial margins, amplified by local media. This is the gold standard for municipal policy-outcome linkage outside the U.S.
- **Difference-in-differences with staggered adoption:** Bolet et al. (2023, APSR) compared Spanish municipalities affected by just-transition agreements to control municipalities, measuring electoral effects of compensatory climate policy. (University of Chicago Press) This design maps directly to studying staggered adoption of social innovations.
- **Dynamic latent variable estimation:** Caughey and Warshaw's (2016, 2022) DGIRT methodology estimates annual policy liberalism scores from sparse binary indicators. If sufficient city-level policy indicators can be assembled, this method generates continuous, dynamic measures of city-level policy orientation linkable to electoral data—the most promising path to operationalizing "urban social innovativeness" as a continuous variable.

The key lesson: **research designs work when they combine a credible identification strategy with granular, linkable data.** The project needs both innovation in measurement (building the city-level policy adoption database) and careful design choices (selecting settings where quasi-random variation exists).

A realistic, staged research design

Given the data landscape, the project should adopt a **three-tier architecture** that is honest about where global comparison works and where it doesn't.

Tier 1 — U.S. deep analysis (primary empirical setting). Construct a database of 50–100 trackable social innovations across the largest 200–500 U.S. cities, recording adoption dates. Candidate innovations include climate emergency declarations, participatory budgeting, inclusionary zoning, community land trusts, guaranteed income pilots, civilian police oversight boards, sanctuary city designations, open data portals, earned sick leave, ban-the-box hiring, plastic bag ordinances, and minimum wage increases above state floors. Sources include CityHealth (12 policies, 75 cities, 2017–present), [City Health](#) NLIHC tenant protections (2021–present), [National Low Income Housing ...](#) ICMA surveys, LocalView transcripts (NLP extraction), and purpose-built coding from municipal websites and ordinance databases. Compute Walker-style innovativeness scores using SPID methodology. [Sage Journals](#) [Wiley Online Library](#) Link to ALGED election data, NCCS civic capacity measures, and Census demographics. Primary designs: staggered-adoption DiD examining incumbent vote share changes after innovation adoption; close-election RD examining whether narrowly elected progressives who innovate face differential backlash; event-study designs around visible policy adoptions. **This tier is feasible within 18–24 months.**

Tier 2 — European comparative analysis (secondary setting). Focus on 3–4 countries with strong municipal data: Germany (comprehensive municipal election database + fiscal data), Spain (municipal elections + just-transition variation + Decidim adoption), Scandinavian countries (NUTS 3 ≈ municipality-level + strong administrative data), and potentially the UK (ward-level elections + social enterprise data). Use EU-NED election data, Eurostat socioeconomic indicators, CDP climate reporting, ESID social innovation project density (geocoded to cities), [PubMed Central](#) and ESEM social enterprise data. When Lo-QoG becomes available, integrate governance quality indicators. [Rj](#) Design: cross-national DiD comparing innovation-electoral outcome relationships across institutional contexts (strong mayor vs. council systems; proportional vs. majoritarian local elections). **This tier is feasible within 24–36 months, contingent on Lo-QoG release.**

Tier 3 — Global climate-focused module (narrowest but most global). Use CDP-ICLEI data (~1,000 cities, 75+ countries) to measure climate action innovativeness—[CDP](#) [CDP](#)a well-defined subset of social innovation. Link to national legislative election results via CLEA constituency data spatially matched to city boundaries, or to mayoral election data where available (Brazil, Germany, U.S.). This enables testing whether climate-innovative cities experience different electoral dynamics, controlling for national trends. **This tier is feasible as a standalone paper within 12–18 months using existing data.**

Tier 4 — Brazil deep dive (strongest non-U.S. within-country option). Brazil offers comprehensive municipal election data (TSE), a long history of participatory budgeting, [Wikipedia](#) random audit natural experiments, and growing What Works Cities infrastructure. A Brazil module studying PB adoption, climate actions, and their electoral consequences would complement the U.S. analysis and test generalizability in a very different institutional context. **Feasible within 18–24 months.**

Honest assessment of what's missing and what's hard

Five challenges deserve candid acknowledgment. First, **the city-level policy adoption database is the binding constraint.** Building it requires extensive hand-coding, web scraping, and administrative data assembly. The LocalView NLP approach can accelerate this for U.S. cities but requires validation against ground-truth policy records. Second, **selection into innovation is severe.** Cities that adopt social innovations differ systematically from those that don't—they are larger, more liberal, have denser civic infrastructure, and face different

problems. Staggered adoption helps but doesn't fully solve selection. Third, **the causal chain is long.** Innovation → implementation → visibility → attribution → electoral response involves multiple mediating steps, any of which can attenuate effects. Arnold's traceability logic suggests that only highly visible, attributable innovations should produce electoral consequences—the project should focus measurement on these. Fourth, **state preemption complicates interpretation.** When states preempt local innovations (Psjblog) (NLC data shows preemption expanding from an average of 3 to 4 policy areas per state, 2019–2024), the "treatment" is censored. (National League of Cities) Preemption must be modeled as a competing risk. Fifth, **the global comparison aspiration must be tempered.** True global comparison at the city level is currently limited to climate policy; broader social innovation comparison is feasible only within countries or within Europe. The project should embrace this honestly rather than overclaiming global scope.

Conclusion

This project sits at a genuinely productive intersection. The political science infrastructure for studying local elections has improved dramatically with the ALGED database. (DNTB) (ResearchGate) The sociology of urban innovation has matured through Brandtner's civic capacity framework. (University of Chicago Press) The methodological toolkit—SPID-style innovativeness scoring, (Zhihu) DGIRT dynamic estimation, text-as-data from LocalView, close-election RD designs—is ready for application at the city level. What's missing is the connective tissue: a structured database of municipal social innovation adoption that can be joined to electoral data. Building that database is the project's most important and most laborious task, but also its most durable contribution. The staged design—U.S. deep analysis first, European comparison second, global climate module throughout, Brazil as a robustness check—matches ambition to data availability. The team should resist the temptation to wait for perfect global data and instead begin with what's achievable, publishing modular findings that progressively expand the scope. The first paper—climate innovation and electoral consequences using CDP + CLEA/ALGED—can be produced within a year and would establish the project's empirical credibility while the larger database is under construction.